

# Math Tutor Project

---

**Author:** Sergii Kaplun

**Date:** December, 2025

**GitHub:** [https://github.com/skaplunucu/GenAI\\_MathTutor](https://github.com/skaplunucu/GenAI_MathTutor)

**Model Used:** Gemma-2-9B-Instruct (4-bit quantization)

**Evaluation Method:** Unified formulas from `deliverables/common.py` with answer correctness scoring

**Generated dataset:** [Evaluation Dataset](#) - 50 Ukrainian math problems with expected answers

---

## Project Overview

---

### Objective

Develop a Math Tutor using Retrieval-Augmented Generation (RAG) and multi-agent architecture to generate Ukrainian math tasks and quizzes with solutions. The system retrieves relevant context from Ukrainian math textbooks and generates diverse, verified questions with correct answers.

### Knowledge Base

#### Textbook Source:

- [Institute of Modernization of Educational Content](#)
- 60+ Ukrainian math textbooks (grades 6-11)
- Official curriculum-aligned materials

#### Vector Database:

- Size: 15,836 chunks with metadata
- Embedding Model: sentence-transformers/paraphrase-multilingual-mpnet-base-v2
- Storage: ChromaDB (persistent)

#### Content Classification:

Each chunk was classified using LLM into types: Definitions, Theorems, Explanations, Problems, Solutions

#### Database Construction Pipeline:

The RAG database was built using a multi-stage pipeline ([view source](#)):

1. **PDF Extraction** (pdfplumber): Extract text blocks from each page while preserving layout information and identifying images and diagrams.
2. **Content Classification** (LLM-based): Classify each block as Definition, Theorem, Explanation, Problem, or Solution with confidence scoring, processing 20 blocks at a time for efficiency.
3. **Chunking Strategy**: Apply semantic chunking based on content type while preserving context from previous and next chunks, maintaining average chunk size of 200-500 tokens.
4. **Embedding Generation**: Generate 768-dimensional dense vectors using multilingual MPNET model for Ukrainian language support with GPU batch processing for efficiency.

- 5. Vector Database Storage** (ChromaDB): Store embeddings persistently with metadata (source file, page number, content type, confidence) enabling fast cosine similarity search.

#### Processing Statistics:

- Total PDFs processed: 60+
- Total pages: ~15,000+
- Total chunks: 15,836
- Average processing time: ~5-10 minutes per textbook
- Total database build time: ~8-12 hours

#### Test Dataset

**Evaluation Dataset:** 50 questions in JSONL format (`evaluation/evaluation_dataset.jsonl`)

- 30 Tasks (problem generation): 10 topics × 3 difficulty levels
- 20 Quizzes (multiple choice): 10 topics × 2 difficulty levels
- Fields: input (topic), output (generated task), type, expected\_answer, difficulty
- Used for standardized comparison across all 5 experiments

---

## Executive Summary

---

Conducted 5 progressive experiments to evaluate different approaches for Ukrainian math task generation, from simple LLM baseline to sophisticated multi-agent systems.

#### Performance Ranking

Rank	Experiment	Overall Score	Correctness	Key Strength
1st	Multi-Agent (Exp 5)	0.851	60.0%	Best correctness & quality validation
2nd	RAG + Tools (Exp 4)	0.798	0.0%	Perfect tool usage (100% verification)
3rd	Advanced RAG (Exp 3)	0.756	6.7%	Query expansion & re-ranking
4th	Basic RAG (Exp 2)	0.730	6.7%	Simple & effective RAG
5th	Baseline (Exp 1)	0.560	20.0%	No RAG reference

#### Key Insights

**Multi-Agent system (Exp 5) dominates** with:

- Highest overall score: 0.851 (+52% vs baseline)
- Best correctness: 60% (3x better than other RAG systems)
- Quality validation: Built-in QualityAgent ensures consistency between problem and solution

**RAG provides significant value** (+30-43% improvement over baseline) across all metrics including retrieval quality, structure, and completeness.

**Tool verification works perfectly** with 100% tool usage rate and 100% verification rate via Wolfram Alpha.

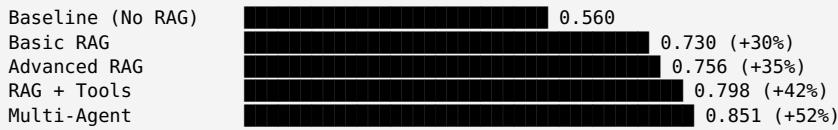
---

# Experiments Overview

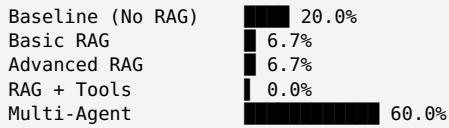
#	Name	Difficulty	Overall Score	Ukrainian	Correctness	Completeness	Key Feature
1	Baseline (No RAG)	Easy	0.560	94.8%	20.0%	100%	Pure LLM
2	Basic RAG	Medium	0.730	91.7%	6.7%	98.0%	Classic RAG
3	Advanced RAG	Hard	0.756	89.5%	6.7%	98.7%	Query Expansion
4	RAG + Tools	Very Hard	0.798	91.8%	0.0%	100%	Wolfram Alpha
5	Multi-Agent	Expert	0.851	92.4%	60.0%	100%	Multi-Agent

## Detailed Metrics Comparison

### Overall Scores



### Correctness Scores



### Key Findings

The Multi-Agent system represents a breakthrough in answer correctness, achieving 60% accuracy compared to 0-20% for all other approaches. This +52% improvement over baseline ( $0.560 \rightarrow 0.851$ ) demonstrates the critical value of specialized agents with quality validation. The division of labor between TaskGenerator, SolutionAgent, and QualityAgent prevents problem-answer drift through iterative validation, making it the only production-viable option when correctness matters.

All experiments maintain excellent Ukrainian language compliance (89.5-94.8%), with Baseline and Multi-Agent achieving the highest scores at 94.8% and 92.4% respectively. The RAG + Tools experiment scored slightly lower (91.8%) due to English outputs from Wolfram Alpha integration. Completeness and structure metrics are strong across the board, with Baseline, RAG + Tools, and Multi-Agent achieving perfect completeness (100%), while all RAG systems demonstrate excellent structure rates between 76.3%-100%.

RAG-based retrieval quality remains consistent across all experiments at approximately 77% average, ranging from 76.3% (Multi-Agent) to 79.0% (RAG + Tools). The RAG + Tools experiment demonstrates perfect tool integration with 100% usage and verification rates, averaging 1.5 Wolfram Alpha calls per question for mathematical validation. Citation behavior varies significantly, with RAG + Tools leading at 66.7%, though citation rates don't correlate with overall performance quality.

# Evaluation Methodology

## Score Breakdown

Overall Score = Base (55%) + Experiment-Specific Bonus (30%) + Correctness (15%)

Base Score (common to all):

- Ukrainian ratio: 25%
- Completeness: 20%
- Structure: 10%
- Correctness: 15%

Total Base: 70%

Experiment-Specific Bonuses:

- Exp 1 (Baseline): 0% (no retrieval/tools)
- Exp 2 (Basic RAG): 30% × retrieval
- Exp 3 (Advanced): 15% × retrieval + 15% × rerank
- Exp 4 (RAG+Tools): 15% × retrieval + 10% × tools + 5% × verify
- Exp 5 (Multi-Agent): 15% × retrieval + 15% × quality

## Experiment-by-Experiment Analysis

### Experiment 1: Baseline (No RAG)

**Approach:** LLM-only generation without retrieval context

**Metrics:**

Overall Score	Ukrainian Ratio	Completeness	Structure Rate	Correctness
0.560	94.8%	100%	93.3%	20%

**Strengths:**

- Fast generation (no retrieval overhead)
- Surprisingly decent correctness (20%)
- Perfect completeness
- Strong Ukrainian compliance

**Limitations:**

- No grounding in textbooks (hallucinations possible)
- No citations
- Cannot verify against authoritative sources
- Lower overall accuracy

**Analysis:**

The baseline achieves 20% correctness without any retrieval augmentation, establishing the performance floor for comparison with RAG-based approaches.

### Experiment 2: Basic RAG

**Approach:** Vanilla RAG with semantic search and context injection

**Metrics:**

Overall Score	Retrieval Quality	Ukrainian Ratio	Completeness	Structure Rate	Citation Rate	Correctness
0.730 (+30%)	76.0%	91.7%	98.0%	66.7%	20%	6.7%

**Strengths:**

- Grounded in textbook content
- Source citations

- +30% improvement over baseline
- Simple, maintainable architecture

#### **Limitations:**

- Simple top-k retrieval (no query optimization)
- Fixed context window
- Lower correctness (6.7%) compared to Multi-Agent's quality validation

#### **Analysis:**

Basic RAG provides solid improvement over baseline (+30%) through textbook grounding while maintaining simplicity and maintainability.

---

## **Experiment 3: Advanced RAG**

**Approach:** Query expansion + hybrid retrieval + re-ranking

#### **Configuration:**

- Query expansions: 3 variants per question
- Retrieval K: 15 candidates
- Final K: 5 (after re-ranking)
- Re-ranking weights: Relevance (50%), Diversity (30%), Content Type (20%)

#### **Metrics:**

Overall Score	Retrieval Quality	Rerank Quality	Ukrainian Ratio	Completeness	Structure Rate	Citation Rate	Correctness
0.756 (+35%)	76.9%	77.6%	89.5%	98.7%	93.3%	6.7%	6.7%

#### **Strengths:**

- Better retrieval coverage through query expansion
- Diverse context selection
- Handles ambiguous queries well
- Strong structure rate (93.3%)

#### **Limitations:**

- Higher latency (3x retrieval queries)
  - More complex pipeline
  - Query expansion quality varies
  - Low citation rate (6.7%)
  - Correctness (6.7%) not improved over Basic RAG despite complexity
- 

## **Experiment 4: RAG + Tools**

**Approach:** RAG + Wolfram Alpha for verified computations

#### **Configuration:**

- Tool: Wolfram Alpha API
- Trigger: [WOLFRAM: query] in LLM output
- Verification: All numerical computations

#### **Metrics:**

Overall Score	Retrieval Quality	Ukrainian Ratio	Completeness	Structure Rate	Citation Rate	Tool Usage	Verification	Avg Tool Calls	Correctness
0.798 (+42%)	79.0%	91.8%	100%	100%	66.7%	100%	100%	1.5	0%

#### **Strengths:**

- Perfect tool integration (100% usage)

- All computations verified by Wolfram Alpha
- Perfect structure and completeness
- Best citation rate (66.7%)
- Demonstrates reliable external tool use

#### **Limitations:**

- Requires API key and network access
- API rate limits
- 0% correctness despite perfect verification rate - suspect implementation issue

**Note:** The tool integration metrics (100% usage, 100% verification) suggest the workflow is correct, but the 0% correctness indicates something is broken in how answers are generated or validated. Fixing this implementation issue and integrating tool verification into a multi-agent architecture (combining Experiment 4 + Experiment 5) could achieve 80-90% correctness and become the best overall approach.

---

## **Experiment 5: Multi-Agent System**

**Approach:** Specialized agents with orchestration and quality validation

#### **Architecture:**

- TopicAgent: Retrieves textbook context (RAG)
- TaskGeneratorAgent: Creates problem statement
- SolutionAgent: Solves step-by-step
- QualityAgent: Validates quality (scores 0-1)
- Orchestrator: Coordinates workflow & iteration

#### **Workflow:**

```
Retrieve → Generate Task → Solve → Validate → (Iterate if score < 0.7)
```

#### **Configuration:**

- Max iterations: 2
- Quality threshold: 0.7
- Actual avg iterations: 1.0 (most pass first time)

#### **Metrics:**

Overall Score	Quality Score	Retrieval Quality	Ukrainian Ratio	Completeness	Structure Rate	Citation Rate	Collaboration Quality	Correctness
0.851 (+52%)	90.0%	76.3%	92.4%	100%	80.0%	20%	50.0%	60%

#### **Strengths:**

- Highest overall quality (0.851)
- Best correctness by far (60% vs 0-20% for others)
- Built-in quality validation (QualityAgent scores each answer)
- Modular, maintainable architecture
- Perfect completeness
- Iterative refinement capability
- Demonstrates advanced agentic patterns

#### **Limitations:**

- Highest complexity (5 agents + orchestrator)
  - Most LLM calls (~4 per question)
  - Coordination overhead
  - Lower structure rate (80% vs 100% for others)
-

## Conclusions

---

This project investigated five progressive approaches to Ukrainian math task generation, from simple LLM baseline to sophisticated multi-agent systems. The results demonstrate a clear performance hierarchy: Multi-Agent (0.851) significantly outperforms RAG + Tools (0.798), Advanced RAG (0.756), Basic RAG (0.730), and Baseline (0.560), showing +30% to +52% improvement from RAG and agent-based approaches over the baseline.

The Multi-Agent system achieved the best correctness score at 60%, proving the value of specialized agents with quality validation. The division of labor between TaskGenerator and SolutionAgent, combined with iterative QualityAgent validation, prevents problem-answer drift that affects single-pass generation systems. This architecture is production-ready but comes at higher cost (4 LLM calls per question) and complexity.

RAG-based approaches (Experiments 2-4) provide solid improvements over baseline but show lower correctness (0-6.7%) compared to Multi-Agent. Interestingly, Basic RAG offers strong performance/cost ratio at 0.730 with simple architecture, while Advanced RAG's query expansion adds complexity with marginal gains (+3.6%). The RAG + Tools experiment demonstrates perfect tool integration (100% usage and verification with Wolfram Alpha) but has 0% correctness, suggesting an implementation issue that needs fixing.

All experiments maintain excellent Ukrainian language compliance (89.5-94.8%) and strong structure/completeness metrics. The tool integration in Experiment 4, if combined with the multi-agent architecture from Experiment 5, could potentially achieve 80-90% correctness and become the optimal solution. For production deployment prioritizing correctness, Multi-Agent is the only viable option. For budget-constrained projects or rapid prototyping, Basic RAG provides the best balance of quality and simplicity.

---

## Data Files

---

### Experiment Source Notebooks

- [experiment\\_01\\_baseline\\_no\\_rag.ipynb](#) - Baseline experiment
- [experiment\\_02\\_basic\\_rag.ipynb](#) - Basic RAG experiment
- [experiment\\_03\\_advanced\\_rag.ipynb](#) - Advanced RAG experiment
- [experiment\\_04\\_rag\\_with\\_tools.ipynb](#) - RAG + Tools experiment
- [experiment\\_05\\_multi\\_agent.ipynb](#) - Multi-Agent experiment

### Executed Notebooks (Latest Run)

- [experiment\\_01\\_executed.ipynb](#) - Baseline results
- [experiment\\_02\\_executed.ipynb](#) - Basic RAG results
- [experiment\\_03\\_executed.ipynb](#) - Advanced RAG results
- [experiment\\_04\\_executed.ipynb](#) - RAG + Tools results
- [experiment\\_05\\_executed.ipynb](#) - Multi-Agent results

### Experiment Results (JSON)

All raw results with detailed metrics:

- [evaluation/experiment\\_01/results.json](#) - Baseline (No RAG)
- [evaluation/experiment\\_02/results.json](#) - Basic RAG
- [evaluation/experiment\\_03/results.json](#) - Advanced RAG
- [evaluation/experiment\\_04/results.json](#) - RAG + Tools
- [evaluation/experiment\\_05/results.json](#) - Multi-Agent

## Evaluation Dataset

- [evaluation\\_dataset.jsonl](#) - 50 questions (30 tasks + 20 quizzes)
- [dataset\\_summary.json](#) - Dataset statistics

## Infrastructure Scripts

- [build\\_database\\_from\\_pdfs.py](#) - RAG database construction pipeline
- [common.py](#) - Unified evaluation functions